

groups, the group keys, the group membership (expressed in terms of binding and ACL entries), as well as the CAT value and version for each group. This section describes how all changes to the Joint Fabric access control configuration are made via the Joint Fabric Datastore: a change is first made to the Datastore where the impacted configuration of individual nodes is marked as pending; the Datastore is then responsible for propagating the change to all impacted nodes and then updating its per-node (and sometimes per-endpoint) pending state to committed.

The Joint Fabric Datastore cluster server SHALL only be accessible on a Node which is acting as the Joint Fabric Anchor Administrator. When not acting as the Joint Fabric Anchor Administrator, the Joint Fabric Datastore cluster SHALL NOT be accessible.

The **Admin** level of access to the Joint Fabric Datastore cluster server SHALL be limited to JF Administrator Nodes identified using the **Administrator CAT**.

**NOTE** Support for Joint Fabric Datastore cluster is provisional.

### 11.24.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

| Revision | Description      |
|----------|------------------|
| 1        | Initial revision |

### 11.24.2. Classification

| Hierarchy | Role    | Scope | PICS Code |
|-----------|---------|-------|-----------|
| Base      | Utility | Node  | JFDS      |

### 11.24.3. Cluster ID

| ID     | Name                                  | Conformance |
|--------|---------------------------------------|-------------|
| 0x0752 | <b>Joint Fabric</b> Datastore Cluster | P           |

### 11.24.4. Usage Requirements

This section describes the required flows for Joint Fabric Administrators to perform common actions related to adding and removing Nodes to a Joint Fabric, and managing Groups and Group Membership on the Joint Fabric.

Many flows performed by a Joint Fabric Administrator require commands to be sent to the Datastore in order to keep the Datastore up-to-date with regard to Nodes and Groups on the Fabric. While a Joint Fabric Administrator SHALL update the Datastore with changes made to Nodes and Groups, it MAY complete these configuration and commissioning tasks when the Datastore is unreachable. If attempts to update the Datastore fail due to temporary conditions (reachability issues, BUSY response error, RESOURCE\_EXHAUSTED error, etc.), the Commissioner SHALL periodically retry the